

Assignment 1

In this assignment you will complete the implementation of a ray caster and parallelize it on the CPU.

- 0.) Download the [skeleton code](#) and generate the build system using `cmake`.¹
- 1.) Complete the implementation of the ray caster. The default scene currently implemented in `Scene` is shown in Fig. 1.
 - i.) Implement `Renderer::render()` that iterates over the pixels in the image, for each pixel generates a ray, and shades the pixel when an intersection point has been found.
 - ii.) Implement `Scene::traceRay()` that iterates over all objects and finds the intersection point closest to the camera.
 - iii.) Implement `Sphere::intersect()` and `Plane::intersect()`.
 - iv.) Implement `Renderer::shade()` that iterates over all lights and calls `Intersection::shade()`.
 - v.) Implement the Phong shading model (as discussed in the tutorial) in `Phong::shade()`.
- 2.) Parallelize the ray caster using C++ threads.
 - i.) Implement `ParallelRenderer::init()` that determines the number of threads as a function of the hardware capabilities.
 - ii.) Implement `ParallelRenderer::render()` that sets up and spawns the given number of threads to render the image. Ensure that the image plane is tiled in an approximately equal manner in the x and y direction.
 - iii.) Implement `ParallelRenderer::renderTile()` that computes a tile of the image.
- 3.) Analysis of execution time.

¹Under Linux and MacOS you can use a package manager (`apt`, `brew`, `port`,...) to install `cmake`. Then type `cmake ./build` on the command line to generate the make file. On Windows `cmake -G "Visual Studio 12" .\build` generates a Visual Studio solution.




Figure 1: Default scene

- i.) Generate a plot with the execution time as a function of the number of threads (see Fig. 2 for an example).
- ii.) Analyse and explain the observed performance in the plot. Suggest (and / or implement) ways to improve the efficiency.

Notes:

- Please finish the implementation until the week of 18/11/2016 and sign up for a time slot to discuss your implementation.
- Please make sure that your code compiles and executes on the machines in the lab or bring your own machine.
- Please email gpu-course@isg.cs.ovgu.de for bug reports (with diff if possible) or comments.




Figure 2: Execution time as a function of the number of threads.

- Plagiarism policy: you can discuss the assignment but everybody has to implement their own solution.